

CS 6240: Assignment 2

Name: Mahesh Bhandarkar
Class: Tuesday 1:35 PM
Homework: HW 2

Source Codes with Highlighted Differences

Custom-Partitioner Code

This code is same for all the variations of the Word Count program we are trying to experiment with.

My “Partitioner” class is as follows:

```
// Output types of Mapper should be same as arguments of Partitioner
public static class customPartitioner
    extends Partitioner<Text, IntWritable>
{
    @Override
    /*
     * Returns m - OneOf['m', 'n', 'o', 'p', 'q']
     * This determines the reduce partition that the word will go to
     * Word starting from 'm' will goto partition 0
     *                               'n' will goto partition 1
     *                               'o' will goto partition 2
     *                               'p' will goto partition 3
     *                               'q' will goto partition 4
     * */
    public int getPartition
        (Text key, IntWritable value, int numPartitions)
    {
        String myKey = key.toString().toLowerCase();
        char firstCharOfKey = myKey.charAt(0);
        return firstCharOfKey - 'm';
    }
}
```

NO-Combiner Source Code

In this case we have modified the Map Method, by adding a filter 'isReal(arg)' That checks if the mapped word is a 'Real' word according to specifications.

The Partitioner , then segregates these words into separate reduce tasks.
'm/M' go to reduce task '

My map class is as follows:

```
public static class Map extends Mapper<LongWritable, Text, Text, IntWritable>
{
    private final static IntWritable one =
        new IntWritable(1);
    private Text word = new Text();
    public void map(LongWritable key, Text value,
        Context context)
        throws IOException, InterruptedException
    {
        String line = value.toString();
        StringTokenizer tokenizer =
            new StringTokenizer(line);
        String strWord;
        char firstLetter;
        while (tokenizer.hasMoreTokens()) {
            word.set(tokenizer.nextToken());
            strWord = word.toString();
            firstLetter = strWord.charAt(0);
            // Check if word is real
            if (this.isReal(firstLetter))
            {
                context.write(word, one);
            }
        }
    }
}

// Checks if the Word is Real or not, by accepting its 1st Character
private boolean isReal(char aFirstCharOfWord)
{
    boolean result =
        Character.toLowerCase(aFirstCharOfWord) == 'm' ||
        Character.toLowerCase(aFirstCharOfWord) == 'n' ||
        Character.toLowerCase(aFirstCharOfWord) == 'o' ||
        Character.toLowerCase(aFirstCharOfWord) == 'p' ||
        Character.toLowerCase(aFirstCharOfWord) == 'q';

    return result;
}
```

```
}  
  
}
```

My Main class for 'NOCombiner' is as follows:

```
public static void main(String[] args) throws Exception {  
    Configuration conf = new Configuration();  
  
    Job job = new Job(conf, "wordcount");  
  
    job.setJarByClass(WordCount.class);  
    // Setting the Partitioner  
    job.setPartitionerClass(customPartitioner.class);  
    job.setOutputKeyClass(Text.class);  
    job.setOutputValueClass(IntWritable.class);  
  
    job.setMapperClass(Map.class);  
    // Setting the Combiner here  
    //job.setCombinerClass(Reduce.class);  
    job.setReducerClass(Reduce.class);  
  
    job.setInputFormatClass(TextInputFormat.class);  
    job.setOutputFormatClass(TextOutputFormat.class);  
  
    FileInputFormat.addInputPath(job, new Path(args[0]));  
    FileOutputFormat.setOutputPath(job, new Path(args[1]));  
  
    job.waitForCompletion(true);  
}
```

Si-Combiner Source Code

In this case, the only change from the former code, is that we introduce a reducer in the Main class, which looks like: “`job.setCombinerClass(Reduce.class);`”

My map class for Si-Combiner is as follows:

The Si-Combiner Map class is same as that of the No-Combiner.

My Main class for ‘Si-Combiner’ is as follows:

```
public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();

    Job job = new Job(conf, "wordcount");

    job.setJarByClass(WordCount.class);
    // Setting the Partitioner
    job.setPartitionerClass(customPartitioner.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);

    job.setMapperClass(Map.class);
    // Setting the Combiner here
    job.setCombinerClass(Reduce.class);
    job.setReducerClass(Reduce.class);

    job.setInputFormatClass(TextInputFormat.class);
    job.setOutputFormatClass(TextOutputFormat.class);

    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));

    job.waitForCompletion(true);
}
```

PerMapTally - Source Code

In this case we introduce a HashMap to work as an in-map combiner. This map gets initialized per-map-call, hence the name PerMapTally.

We need to comment out the combiner from the main class, in order to see the hashmap work as an in-mapper combiner.

My map class for PerMapTally is as follows:

```
public static class Map extends Mapper<LongWritable, Text, Text, IntWritable>
{
    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();
    private IntWritable iwCount = new IntWritable();

    // localMap store count for each map call
    private HashMap<String, Integer> localMap;

    public void
    map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {

        // Local map
        localMap = new HashMap<String, Integer>();

        String line = value.toString();
        StringTokenizer tokenizer =
            new StringTokenizer(line);
        String strWord;
        char firstLetter;
        while (tokenizer.hasMoreTokens())
        {
            word.set(tokenizer.nextToken());
            strWord = word.toString();
            firstLetter = strWord.charAt(0);
            // Check if word is real
            if (this.isReal(firstLetter))
            {
                if (!localMap.containsKey(strWord))
                {
                    localMap.put(strWord,1);
                }
                else
                {
                    localMap.put(strWord, localMap.get(strWord) + 1);
                }
            }
        }

        // EMIT to Context from Local Map
        for (String wrd : localMap.keySet())
```

```

    {
        word.set(wrd);
        iwCount.set(localMap.get(wrd));
        context.write(word, iwCount);
    }
}

// Checks if the Word is Real or not, by accepting its 1st Character
private boolean isReal(char aFirstCharOfWord)
{
    boolean result =
        Character.toLowerCase(aFirstCharOfWord) == 'm' ||
        Character.toLowerCase(aFirstCharOfWord) == 'n' ||
        Character.toLowerCase(aFirstCharOfWord) == 'o' ||
        Character.toLowerCase(aFirstCharOfWord) == 'p' ||
        Character.toLowerCase(aFirstCharOfWord) == 'q';

    return result;
}

}.

```

My Main class for 'PerMapTally' is as follows:

```

public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();

    Job job = new Job(conf, "wordcount");

    job.setJarByClass(WordCount.class);
    // Setting the Partitioner
    job.setPartitionerClass(customPartitioner.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);

    job.setMapperClass(Map.class);
    // Setting the Combiner here
    // job.setCombinerClass(Reduce.class);
    job.setReducerClass(Reduce.class);

    job.setInputFormatClass(TextInputFormat.class);
    job.setOutputFormatClass(TextOutputFormat.class);

    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));

    job.waitForCompletion(true);
}

```

PerTaskTally - Source Code

In this case we use the previous Hashmap, but only to initialize it at the setup(), and Emit on cleanup().

My map class for PerMapTally is as follows:

```
public static class Map extends Mapper<LongWritable, Text, Text, IntWritable>
{
    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();
    private IntWritable iwCount = new IntWritable();

    // localMap store count for each map call
    private HashMap<String, Integer> localMap;

    public void setup(Context context) throws IOException,
    InterruptedException
    {
        localMap = new HashMap<String, Integer>();
    }

    public void map(LongWritable key, Text value, Context context) throws
    IOException, InterruptedException {

        String line = value.toString();
        StringTokenizer tokenizer = new StringTokenizer(line);
        String strWord;
        char firstLetter;
        while (tokenizer.hasMoreTokens())
        {
            word.set(tokenizer.nextToken());
            strWord = word.toString();
            firstLetter = strWord.charAt(0);
            // Check if word is real
            if (this.isReal(firstLetter))
            {
                if (!localMap.containsKey(strWord))
                {
                    localMap.put(strWord,1);
                }
                else
                {
                    localMap.put(strWord, localMap.get(strWord) + 1);
                }
            }
        }
    }
}
```

```

// Checks if the Word is Real or not, by accepting its 1st Character
private boolean isReal(char aFirstCharOfWord)
{
    boolean result =
        Character.toLowerCase(aFirstCharOfWord) == 'm' ||
        Character.toLowerCase(aFirstCharOfWord) == 'n' ||
        Character.toLowerCase(aFirstCharOfWord) == 'o' ||
        Character.toLowerCase(aFirstCharOfWord) == 'p' ||
        Character.toLowerCase(aFirstCharOfWord) == 'q';

    return result;
}

// EMITTS to the Context
protected void cleanup (Context context)
    throws IOException, InterruptedException
{
    // EMITT to Context for each line in main text
    for (String wrd : localMap.keySet())
    {
        word.set(wrd);
        iwCount.set(localMap.get(wrd));
        //emitt
        context.write(word, iwCount);
    }
}
}

```

My Main class for 'PerMapTally' is as follows:

Same as main class for the PerMapTally, as we need the partitioner and nothing but the Hashmap that works as a combiner upon each map task..

```

public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();

    Job job = new Job(conf, "wordcount");

    job.setJarByClass(WordCount.class);
    // Setting the Partitioner
    job.setPartitionerClass(customPartitioner.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);

    job.setMapperClass(Map.class);
    // Setting the Combiner here
}

```



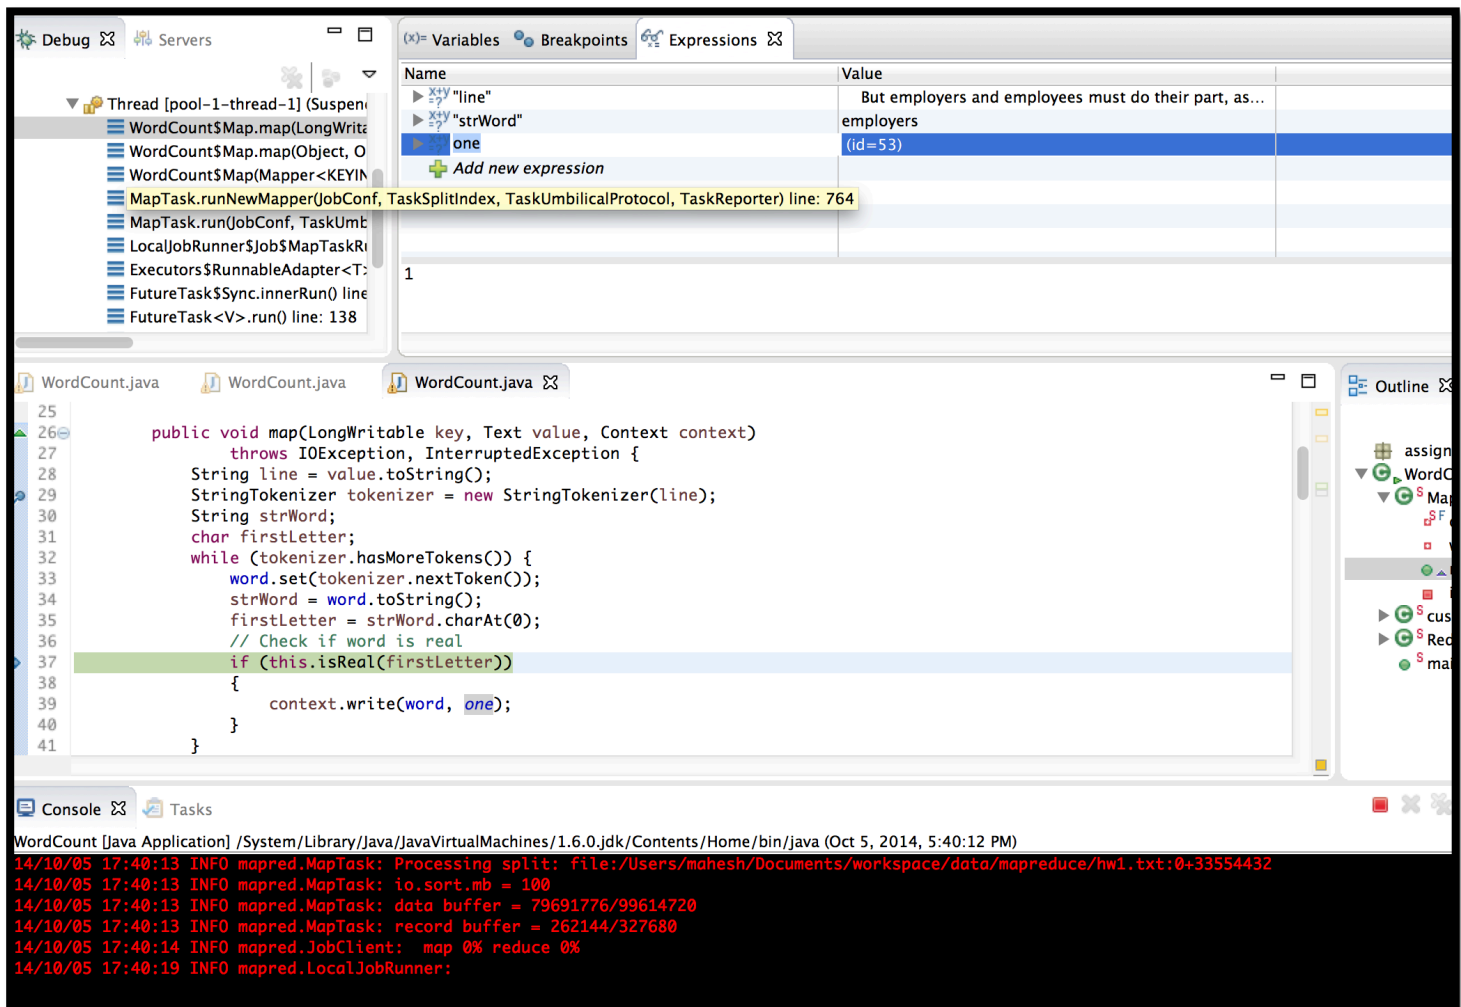
```
// job.setCombinerClass(Reduce.class);
job.setReducerClass(Reduce.class);

job.setInputFormatClass(TextInputFormat.class);
job.setOutputFormatClass(TextOutputFormat.class);

FileInputFormat.addInputPath(job, new Path(args[0]));
FileOutputFormat.setOutputPath(job, new Path(args[1]));

job.waitForCompletion(true);
```

Explanation – Considering the Debug perspective



In the above snapshot, you can see a debug perspective. This is the debug perspective I got, when I was single stepping the No-Combiner program

In the Expressions tab, you can see Name-Value Pair,
Now when the word count program is run, and we single step through the program, I paused at a step to analyze the situation.

The '**line**' is the variable that stores the line from the input text, for each Map-call,
The '**line**' has the following value:

" But employers and employees must do their part, as well, as they are"

'**one**' is the variable with the value **1**, as we count 1 for each word in the map.

'**strWord**' is a variable that has the string form of the "Text" from the incoming input.

strWord in this case has a string '**employers**' stored in it, which according to the specifications is NOT a 'real' word, and would be ignored by our map class.

Performance Analysis:

Analyzing the performance of the 4 variations of word count that are run in 2 configurations. We run these variations twice on the given configurations.

Configurations are as follows:

Configuration 1: 6 small machines (1 master, 5 workers)
Configuration 2: 11 small machines (1 master, 10 workers)

Table of Analysis:

Variations of WordCount	Configuration 1		Configuration 2	
	1 st Run (seconds)	2 nd Run (seconds)	1 st Run (seconds)	2 nd Run (seconds)
NoCombiner	341	339	296	272
SiCombiner	322	301	233	219
PerMapTally	328	371	291	261
PerTaskTally	252	265	196	192

These values were taken from the **controller** log file that was generated after the programs ran on AWS.

Do you believe Combiner was called in the Si Combiner?

Yes.

For an instance I analyzed the Syslog, when I ran the NoCombiner and SiCombiner on the 1-Master-10-Worker configuration.

I analyzed the syslog files for both program-runs, and found following numbers that convince me that the combiner did run in the SiCombiner.

(Taken from 2nd configuration : 10 Worker, 1 Master 2nd Run)

Syslog Entry: NoCombiner

Combine input records=**0**

Combine output records=**0**

Syslog Entry: SiCombiner

Combine input records=**42989997**

Combine output records=**166275**

What difference did the use of Combiner make in the Si Combiner, as compared to the NoCombiner?

There is sure a good amount of difference, the use of a combiner makes.

Analyzing the syslogs for both **SiCombiner** and **NoCombiner** I found that:
(Taken from 2nd configuration : 10 Worker, 1 Master 2nd Run)

Syslog entry : NoCombiner

Reduce input records=**42842400**

Syslog entry : SiCombiner

Reduce input records=**18678**

Clearly we can see a significant change in the input size the Reducer is receiving. Apparently the **reducer** in **SiCombiner** gets **significantly less input**, as compared to **NoCombiner**.

Nevertheless, we also have already analyzed the time taken for **NoCombiner** and **SiCombiner**, to run on AWS, and SiCombiner clearly runs faster than NoCombiner.

Variations of <u>WordCount</u>	Configuration 1		Configuration 2	
	1 st Run (seconds)	2 nd Run (seconds)	1 st Run (seconds)	2 nd Run (seconds)
NoCombiner	341	339	296	272
SiCombiner	322	301	233	219

Was the local aggregation effective in PerMapTally compared to NoCombiner?

Variations of <u>WordCount</u>	Configuration 1		Configuration 2	
	1 st Run (seconds)	2 nd Run (seconds)	1 st Run (seconds)	2 nd Run (seconds)
NoCombiner	341	339	296	272
PerMapTally	328	371	291	261

I believe most of the times, PerMapTally, should be faster than NOCombiner. The numbers in the above table make me believe that.

But in the 2nd run for Configuration 1, the NoCombiner seemed to be faster then PerMapTally.

To conclude, I would say that PerMapTally doesn't make a significant difference by employing local aggregation.

Also to bolster the statement, the syslog for both implementations on 2nd config says, that

PerMapTally's map emits: Map output bytes=396549900

Whereas, NoCombiner's map emits: Map output bytes=412253400

What differences do you see between PerMapTally and PerTaskTally? Try to explain the reason.

The following are the entries I analyzed from the Syslogs of the respective variations of Word Count :

(Taken from 2nd configuration: 10 Worker, 1 Master 2nd Run)

PerTaskTally

FILE_BYTES_READ=97223

FILE_BYTES_WRITTEN=1277758

Map output materialized bytes=191082

Spilled Records=37356

PerMapTally

FILE_BYTES_READ=74564534

FILE_BYTES_WRITTEN=83605195

Map output materialized bytes=27142199

Spilled Records=122598900

Clearly you could see that, the File-bytes Read and written, the Spilled Records are far less in PerTaskTally, than in PerMapTally

Also, from the *Map Output Materialized Bytes*, you could observe the fact that there are far less number of bytes written on to the disk in PerTaskTally, than in PerMapTally

Nevertheless, the *PerTaskTally* runs faster than *PerMapTask* as you can see in the following Table:

Variations of <u>WordCount</u>	Configuration 1		Configuration 2	
	1 st Run (seconds)	2 nd Run (seconds)	1 st Run (seconds)	2 nd Run (seconds)
PerMapTally	328	371	291	261
PerTaskTally	252	265	196	192

The main reason why the map in **PerTaskTally**, outputs lesser byte of data, is because the mapper gets dose some computation on its end, which saves the reducer from doing more work.

The mapper of **PerTaskTally** maps the words, and counts the **repeated occurrences** of the words, thus delivering a word and its count in that particular **map task**, to the reducer.

Wheras in the **PerMapTaly**, the word mapping is done for every line, as we traverse through the massive text, this is much costlier than the way PerTaskTally does it. Thus it takes more time.

Also, **PerMapTaly** maps the words, and counts the **repeated occurrences** of the words, thus delivering a word and its count in that particular **map call**, to the reducer. Now some times this may take same time as NoCombiner, considering a case where there are no repeated words in a line.

Moreover, since the Reducer in the ***PerTaskTally*** has lesser records to process, is faster than ***PerMapTally***.

Which one is better: SiCombiner or PerTaskTally? Briefly justify your answer.

Variations of <u>WordCount</u>	Configuration 1		Configuration 2	
	1 st Run (seconds)	2 nd Run (seconds)	1 st Run (seconds)	2 nd Run (seconds)
SiCombiner	322	301	233	219
PerTaskTally	252	265	196	192

Considering the above table, we see that PerTaskTally is a clear winner as far as Speed goes.

Now let us consider the entries from Syslog, for each of them:

SiCombiner:

```
FILE_BYTES_READ=4416070
FILE_BYTES_WRITTEN=2708373
Map output bytes=412253400
Map output records=42842400
Spilled Records=184953
```

PerTaskTally:

```
FILE_BYTES_READ=97223
FILE_BYTES_WRITTEN=1277758
Map output bytes=229702
Map output records=18678
Spilled Records=37356
```

We observe that the ***PerTaskTally*** has **lower** File Bytes Read, File Bytes Written, Map Output Bytes, Map Output Records and Split Records.
Hence I would regard ***PerTaskTally*** better than ***SiCombiner***.

Comparing the results for configuration 1 & 2, do you believe this MapReduce program scales well to larger cluster? Briefly justify your answer.

While we observe a drastic change in the execution times of these MapReduce programs when on configuration1 versus vonfiguration2, to check for the scalability of this MapReduce program on to a larger cluster, (i.e, from configuration 1 to configuration 2), it is not enough to have the same data set.

As we double the Worker machines, we would also need to double the Data set, as an input to the cluster.

We would require comparing an execution of a MapReduce program, with double data.

With the information we poses:

Variations of <u>WordCount</u>	Configuration 1		Configuration 2	
	1 st Run (seconds)	2 nd Run (seconds)	1 st Run (seconds)	2 nd Run (seconds)
NoCombiner	341	339	296	272
SiCombiner	322	301	233	219
PerMapTally	328	371	291	261
PerTaskTally	252	265	196	192

We can say that, on a larger Cluster, we get a reduced time of execution.

This observation in a way suggests the fact that on increasing the input to double, we still would achieve an approximate similar time of execution, if not same. Hence it would be safe to infer that the MapReduce program scales well.